



Predicting user features from Social Media data

Presented by User 08: The Maze

François Mercier 20167322

Nicolas Sauthier 932337

Zicong Mo 20141760

Andrew Kristensen 20119706

Yifan(Andy) Bai 20153885



What is the problem?



Problem

We have user data from various social media accounts, and we'd like to be able to predict:

- Age
- Gender
- 5 personality traits

We have access to :

- Oxford facial features
- LIWC and NRC features from text posts
- Relation data from liked pages

**How can we use this data to
achieve the goal?**



Page Relations

Drew



Age and Gender

Nicolas and Andy



Personality Predictions

Zicong and François



Early Approach: Page Relations

“Similar people like similar things”

Our earliest idea for estimating user values was utilizing the relations data


Graph structure might be exploitable

Easy and fast to get working for comparisons against the baselines

Approaches: Page Relations

General idea: If we know which kinds of users have liked which pages, we can predict a new user's values based on which pages they have liked


For each user, we hold onto their id and all the pages that they've liked



Approaches: Page Relations

General idea: If we know which kinds of users have liked which pages, we can predict a new user's values based on which pages they have liked

For each page, we hold onto the pageid and all the users who like them

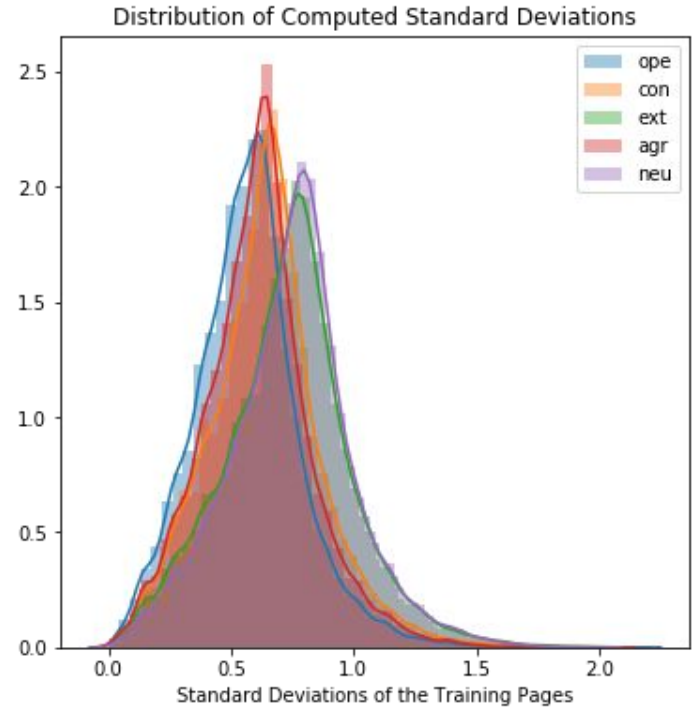


Page Relations - Regression

We can predict the values for the personality traits (ope, con, etc) by taking a weighted average from all the pages the user likes

Why a weighted average?

The more we trust an estimate, the more credit we should give to it



Trait	OPE	CON	EXT	AGR	NEU
Mean	0.5683	0.6386	0.7306	0.6093	0.7469



Page Relations - Regression

How did this go?

Trait	ope	con	ext	agr	neu
Model RMSE	0.613	0.7086	0.7973	0.6565	0.7896
Baseline	0.652	0.798	0.788	0.665	0.734
Difference	-0.039	-0.0894	0.0093	-0.0085	0.0556

Ok, but not great



Page Relations - Classification

We approached the problem of classification for the ages and genders in a binned regression manner

Just like the personalities, we would take a weighted mean

Bin the output value given some thresholds (.5 for the gender, and age limits for the ages)



Page Relations - Classification

We approached the problem of classification for the ages and genders in a binned regression manner

Just like the personalities, we would take a weighted mean

Bin the output value given some thresholds (.5 for the gender, and age limits for the ages)

...This didn't go great




Page Relations - Where is it?

The approaches that we took in the next parts of this presentation out performed it, so it has gone the way of the dodo

We knew this would be the case, since we needed something to give ourselves a slightly improved baseline that we could then beat again.

Since we never implemented any of the Node2Vec work, the model lacked sophistication

Gender Prediction




Data source: all three

1. Face data:
 - a. Most facing or mean imputation
 - b. Correction for yaw and pitch for “front face”
 - c. Distance from nose, resized for eye_dist = 1
 - d. Removing of very highly correlated (>0.99)**TOTAL 25 features**
2. NRC/LIWC data
 - a. Removing of very highly correlated (>0.99)**TOTAL 88 features**
3. Relational data
 - a. Co-occurrence matrix (same userID for two likeids), first 10'000 likeid
 - b. Reduced with SVD to 15 dimensions.
 - c. Average of all user's likeid. Imputations with 0**TOTAL 15 features**

Algorithm: Ada Boost

- Grid search for hyperparameters

Gender Prediction - Results




Age Prediction



- Naive Bayes' wasn't successful with facial data only
- Tried ensemble method
 - With text data (nrc + liwc after feature selection):
 - No resampling: Gradient Boosting
 - Resampling: Random Forest/Extra Tree
 - With facial + text data :
 - AdaBoost



Training Data

- Encode based on age groups: xx-24: 0, 25-34: 1, 35-49: 2, 50+: 3
- Data highly imbalanced



Feature Engineering with Text Data

- Merge NRC and LIWC by user ID after analyzing features, rejecting unuseful ones (e.g. highly correlated/skewed)





Feature Engineering with Text Data

- Both have many features with high number of zeros
- Rejected variables (LIWC):

`'Comma': Highly correlated with 'AllPct' - 0.9408211702`

`'Funct': Highly correlated with 'Dic' - 0.9371835644`

`'QMark': Highly correlated with 'Comma' - 0.9437018647`



Gradient Boost

- Baseline setup results (random state = 42), 5-fold CV: 0.611
- Random Search: 0.615
- Some over-fitting exist, public test accuracy: 0.612 - 0.616



Resampling Efforts

- Training data highly imbalanced, consider resampling
- Extra Trees and Random Forest
- RandomOverSampler:
 - 5-fold CV, train/test split = 0.8/0.2
 - Random Forest = 0.942
 - Extra Trees = 0.961
- But failed miserably with public test (<0.6), due to overfitting in training



AdaBoost



- Adopted Nicolas' feature set
- Grid Search: acc = 0.657 with public test



Personality Predictions

Exploratory Data Analysis

- No orphans
- No missing data
- Skewed distribution for labels and **correlated**






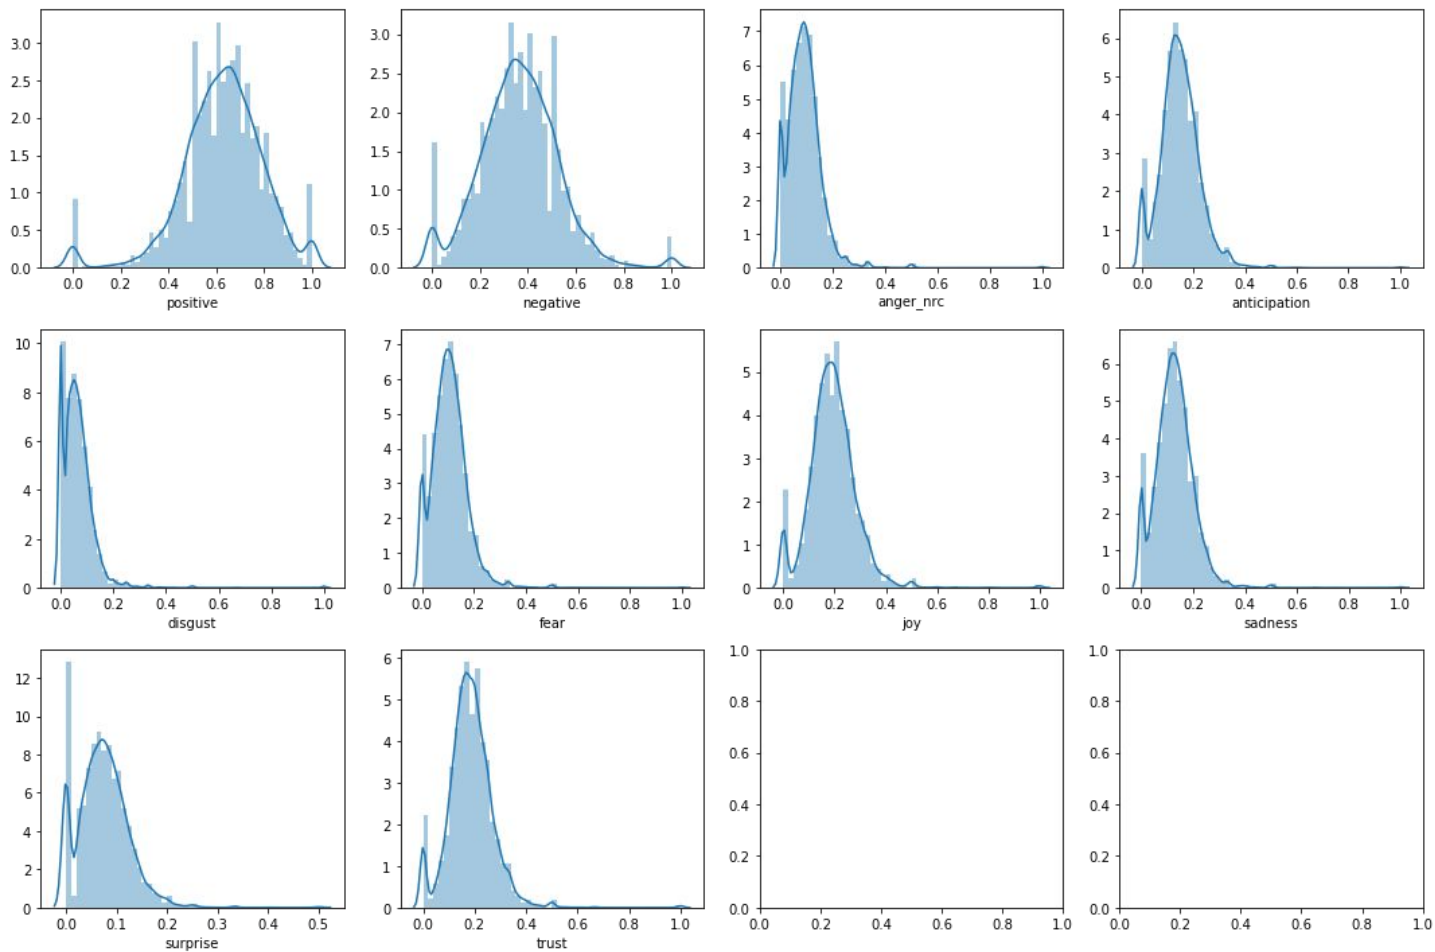
Personality Predictions

Exploratory Data Analysis

- No orphans
- No missing data
- Skewed distribution for labels **and correlated**
- NRCC features
 - Two group: Positive vs negative ones
 - Strong correlation
 - **Presence of outliers (missing data?)**
-



NRCC features distributions








Personality Predictions

Exploratory Data Analysis

- No orphans
- No missing data
- Skewed distribution for labels **and correlated**
- NRCC features
 - Two group: Positive vs negative ones
 - Strong correlation
 - **Presence of outliers (missing data?)**
- LIWC features
 - Groups of features with strong correlations
 - **Presence of outliers (missing data?)**
 -



LIWC features distributions





Personality Predictions

Exploratory Data Analysis

- ⇒ With skewed labels, focus on decision tree family model and ensemble models using it
- ⇒ Focus on removing features to simplify models
- ⇒ Presence of outliers (replacing by median, mean)

How our models worked out

Personality Predictions

Model	Training errors (80%)	Validation errors (20%)	Comments
Mean baseline	Open RMSE 0.62775 Neurotic RMSE 0.7904 Extrovert RMSE 0.8077 Agreeable RMSE 0.6605 Conscientious RMSE 0.7169	Open RMSE 0.6474 Neurotic RMSE 0.8045 Extrovert RMSE 0.8146 Agreeable RMSE 0.6651 Conscientious RMSE 0.7299	
Decision tree	Open RMSE 0.6299 Neurotic RMSE 0.7819 Extrovert RMSE 0.799 Agreeable RMSE 0.6487 Conscientious RMSE 0.7078	Open RMSE 0.6203 Neurotic RMSE 0.8001 Extrovert RMSE 0.8170 Agreeable RMSE 0.6688 Conscientious RMSE 0.7119	91 features Overfitting
Decision tree with L1 features selection	Open RMSE 0.6284 Neurotic RMSE 0.7917 Extrovert RMSE 0.8030 Agreeable RMSE 0.6618 Conscientious RMSE 0.7161	Open RMSE 0.6175 Neurotic RMSE 0.7947 Extrovert RMSE 0.8053 Agreeable RMSE 0.6417 Conscientious RMSE 0.7143	6 features only No overfitting Better performance than with all features Not enough to beat baselines

Personality Predictions

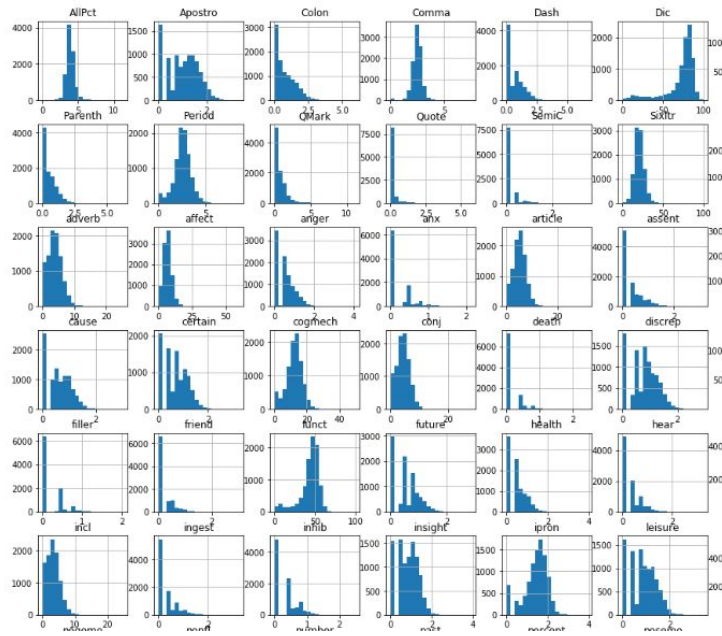

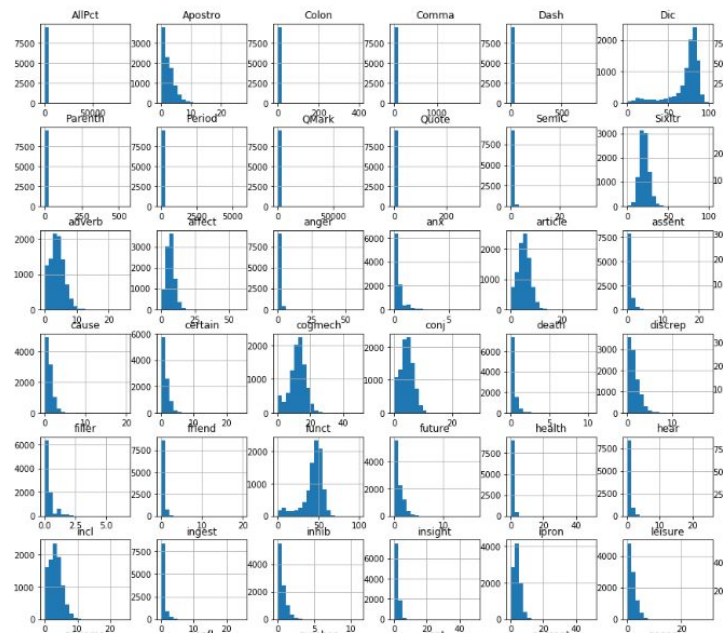
Model	Training errors (80%)	Validation errors (20%)	Comments
Random Forest	Open RMSE 0.6193 Neurotic RMSE 0.7642 Extrovert RMSE 0.7794 Agreeable RMSE 0.6404 Conscientious RMSE 0.6917	Open RMSE 0.6200 Neurotic RMSE 0.7966 Extrovert RMSE 0.8081 Agreeable RMSE 0.6543 Conscientious RMSE 0.7065	All features (as sampling features doesn't help) OOB R^2 score: 2.3% (very poor)
Gradient Boosting	Open RMSE 0.6284 Neurotic RMSE 0.7917 Extrovert RMSE 0.8030 Agreeable RMSE 0.6618 Conscientious RMSE 0.7161	Open RMSE 0.6175 Neurotic RMSE 0.7947 Extrovert RMSE 0.8053 Agreeable RMSE 0.6417 Conscientious RMSE 0.7143	All features (as sampling features doesn't help) No specific management for outliers 1 model per score Using some scores to predict others doesn't help (fusion approach) Manual hyperparameters selection Best model on valid set Bootstrap RMSE to get confidence interval on RMSE (30 tries) Beat the baselines slightly

Personality Predictions



Gradient Boosting model confidence interval Using 80/20 split	Error means On valid set	Low (2 * std dev)	High (2 * std dev)	Z test 1 tailed test	p-value for beating the baseline
opn_rmse	0.620502	0.597431	0.643572	-14.955	0.00%
neu_rmse	0.792591	0.775389	0.809794	-3.444	0.03%
ext_rmse	0.810185	0.790275	0.830096	12.206	100%
agr_rmse	0.662342	0.639646	0.685039	-1.282	9.98%
con_rmse	0.706747	0.684471	0.729023	-13.401	0.00%

Deep Learning







Train on 7600 samples, validate on 1900 samples

Epoch 1/35

7600/7600 [=====] - 3s 458us/step - loss: 1.6652 - val_loss: 0.8075

Epoch 2/35

7600/7600 [=====] - 1s 154us/step - loss: 0.7577 - val_loss: 0.6665



```
model = Sequential()
model.add(Dense(364, input_dim=91, activation="relu"))
model.add(Dropout(0.2))
model.add(Dense(1000, activation="relu"))
model.add(Dense(1000, activation="relu"))
model.add(Dense(500, activation="relu"))
model.add(Dense(250, activation="relu"))
model.add(Dense(100, activation="relu"))
model.add(Dense(20, activation="relu"))
model.add(Dense(5, activation='linear'))
model.compile(loss="mean_squared_error", optimizer="Nadam")
```

```
model = Sequential()
model.add(Dense(40, input_dim=91, activation="relu"))
model.add(Dropout(0.2))
model.add(Dense(300, activation="relu"))
model.add(Dense(300, activation="relu"))
model.add(Dense(300, activation="relu"))
model.add(Dense(100, activation="relu"))
model.add(Dense(5, activation="linear"))
model.compile(loss="mean_squared_error", optimizer="Nadam")
self._model = model
```

```
model = Sequential()
model.add(Dense(500, input_dim=91, activation="relu"))
model.add(Dropout(0.2))
model.add(Dense(1000, activation="relu"))
model.add(Dense(1000, activation="relu"))
model.add(Dense(500, activation="relu"))
model.add(Dense(205, activation="relu"))
model.add(Dense(125, activation="relu"))
model.add(Dense(50, activation="relu"))
model.add(Dense(25, activation="relu"))
model.add(Dense(25, activation="relu"))
model.add(Dense(5, activation="linear"))
model.compile(loss="mean_squared_error", optimizer="Nadam")
self._model = model
```


```
model = Sequential()
model.add(Dense(20, input_dim=91, activation="relu"))
model.add(Dropout(0.2))
model.add(Dense(40, activation="relu"))
model.add(Dense(40, activation="relu"))
model.add(Dense(10, activation="relu"))
model.add(Dense(5, activation="linear"))
model.compile(loss="mean_squared_error", optimizer="Nadam")
```

Personality Predictions

```
model = Sequential()
model.add(Dense(200, input_dim=91, activation="linear"))
model.add(Dropout(0.1))
model.add(Dense(200, activation="linear"))
model.add(Dropout(0.1))
model.add(Dense(100, activation="linear"))
model.add(Dropout(0.1))
model.add(Dense(50, activation="linear"))
model.add(Dropout(0.1))
model.add(Dense(20, activation="linear"))
model.add(Dropout(0.1))
model.add(Dense(5, activation='linear'))
model.compile(loss="mean_squared_error", optimizer="Nadam")
```

NN model confidence interval Using 80/20 split	Error means On valid set	Low (2 * std dev)	High (2 * std dev)	Z test 1 tailed test	p-value for beating the baseline
opn_rmse	0,6198	0.6053	0.6428	-15.40	0%
neu_rmse	0,7907	0.7726	0.8231	-4.909	0.029%
ext_rmse	0,8041	0.7819	0.8240	7.256	100%
agr_rmse	0,6578	0.6360	0.6721	-3.853	0.006%
con_rmse	0,7068	0.6839	0.7225	-20.106	0%

Ensemble Model



Personality
Predictions



Personality
Predictions



50%

50%

Equal weight

Personality
Predictions




Best

Personality
Predictions



Personality
Predictions



Personality
Predictions




Pairwise linear regression
Without intersection

1 2
3 4

Personality
Predictions



Personality
Predictions



Personality
Predictions




Another NN

Personality
Predictions



Personality
Predictions



Personality
Predictions



Multivariate Linear Regression

Personality Predictions

Ensemble Model confidence interval Using 80/20 split	Error means On valid set	Low (2 * std dev)	High (2 * std dev)	Z test 1 tailed test	p-value for beating the baseline
opn_rmse	0.62069 ↑	0.599249 ↓	0.642132 ↓	-15.99 ↓	0%
neu_rmse	0.789116 ↓	0.770993 ↓	0.807239 ↓	-5.369 ↓	0%
ext_rmse	0.798725 ↓	0.774492 ↓	0.822958 ↓	4.8482 ↓	100%
agr_rmse	0.655121 ↓	0.637848 ↑	0.672393 ↑	-6.265 ↓	0%
con_rmse	0.703591 ↓	0.682387 ↓	0.724796 ↓	-15.70 ↓	0%