

Classifying Ultrasonic Mouse Vocalizations using a TensorFlow Network

Drew Kristensen
Dr Adam A Smith

Why classify mouse vocalizations?

What are the characteristics of an autistic mouse?

- It's difficult to create a mouse model of mental diseases

Decoding mouse ultrasonic vocalizations (USVs) might provide insight into their affective state

Potentially helpful in modeling mental disorders

Current methods rely on human operators to label where calls occur and don't occur - time intensive and we have better things to do


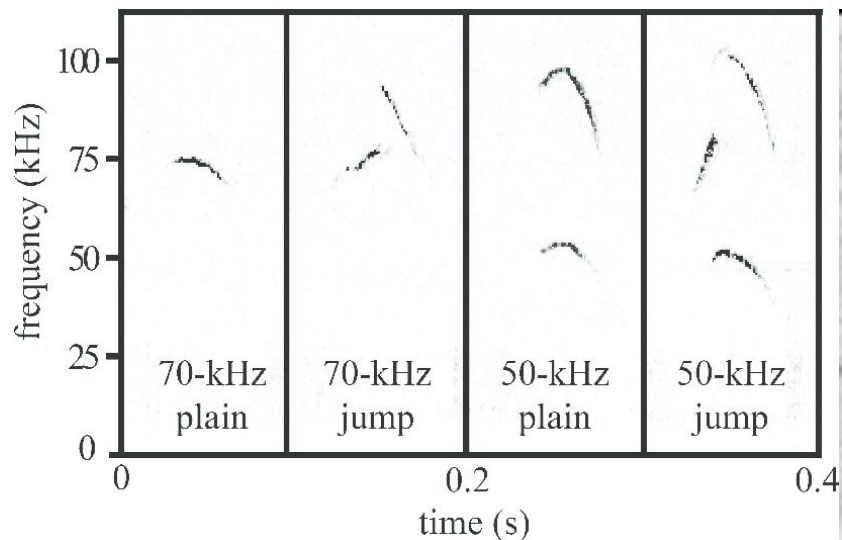
What is the goal?

Classify USVs from recordings using a neural network from the tensor flow framework

Mouse Vocalizations

Often emitted between 30 kHz and 100 kHz

Human hearing goes up to 20 kHz under perfect conditions (hence the 'ultrasonic')

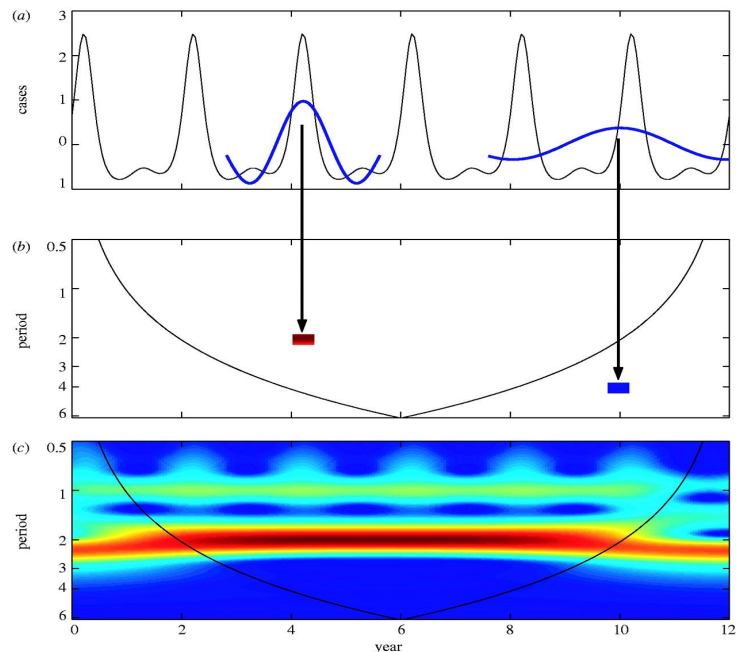



Scalograms

We hear in log-scale, so how can we make our data reflect that?

Scalograms are visual representations of the audio data

Uses wavelet transforms instead of Fourier Series





Neural Network

All the rage in Machine Learning

Useful for classification problems


An attempt to model the abilities of human pattern detection



TensorFlow

Google's open source framework for creating computational data flow graphs

Extremely powerful




Google's LeNet (A really really complex data flow graph)

The Identification Process

Start with .wav files, sampled at 250 kHz


Find the wavelet coefficients between 30 kHz and 120 kHz

For training, annotate the scalogram




The Identification Process

Take the coefficients
around the point we
want to classify and
flatten them into a
vector



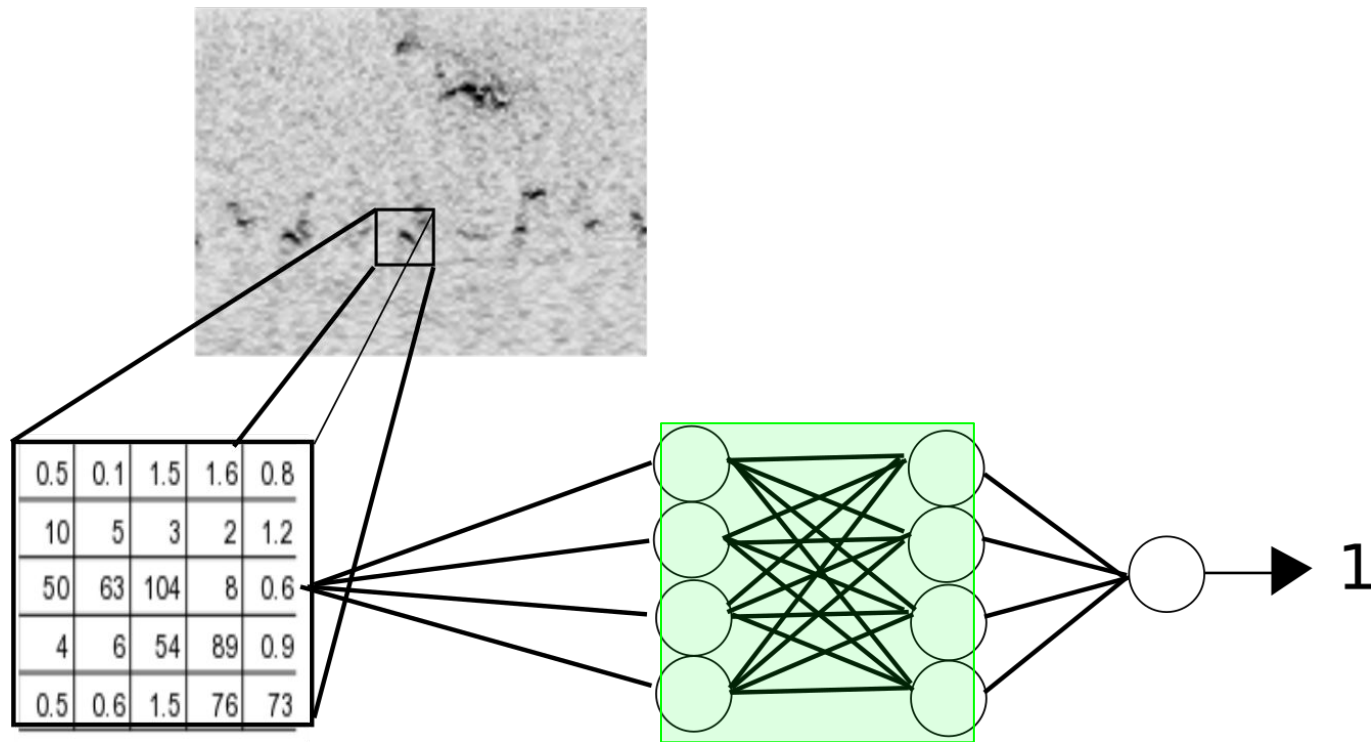
The Identification Process

Feed the input
vector into the first
layer of the neural
network




The Identification Process

The first layer
outputs their
weighted sums to
the second layer




The Classification Process

The second layer
outputs their values
to the final neuron -
the output layer




The Identification Process

Since the input was centered around a call, the network outputs a value of 1




Current Results

Input Scalogram




Elements with output of at least 1



Challenges


Need for human labeled data - not always easy to annotate*



What is a call and what isn't?

Future Directions

Classify different types of calls - not just their presence



Questions?